

Spring 2024: CS5720
Neural Networks and Deep Learning - ICP-7
Yasaswini Majety (700747747)

Github Link: https://github.com/yasaswini8777/Neural_ICP_7

Use Case Description:

LeNet5, AlexNet, Vgg16, Vgg19

1. Training the model
2. Evaluating the model

Programming elements:

1. About CNN
2. Hyperparameters of CNN
3. Image classification with CNN

In class programming:

1. Follow the instruction below and then report how the performance changed.(apply all at once)
 - Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Flatten layer.
 - Dropout layer at 20%.
 - Fully connected layer with 1024 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected layer with 512 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected output layer with 10 units and a Softmax activation function

Did the performance change?
2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.
3. Visualize Loss and Accuracy using the history object

```
Content | Bb 15492383 | colab.google | 700747747_ICP_7.ipynb - C | +
colab.research.google.com/drive/1-o0r9339M5zfilyPyxUrOrJLvwqDZMzO?authuser=0#scrollTo=kt0h1hhprovL
700747747_ICP_7.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
# Simple CNN model for CIFAR-10
import numpy
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
from keras.layers import Flatten
from keras.constraints import maxnorm
from keras.optimizers import SGD
from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.utils import to_categorical
#from keras import backend as K
#K.set_image_dim_ordering('th')

# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
# normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train = X_train / 255.0
X_test = X_test / 255.0
# one hot encode outputs
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
num_classes = y_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
# Compile model
epochs = 5
lr = 0.01
decay = lr/epochs
sgd = SGD(lr=lr)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)
# Final evaluation of the model
20m 46s completed at 9:21 PM
```

```
Content | Bb 15492383 | 700747747_ICP_7.ipynb - C | Neural_ICP_7_700747747 - C | +
colab.research.google.com/drive/1-o0r9339M5zfilyPyxUrOrJLvwqDZMzO?authuser=0#scrollTo=kt0h1hhprovL
700747747_ICP_7.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
sgd = SGD(lr=lr)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 2s 0us/step
WARNING:absl: 'lr' is deprecated in Keras optimizer, please use 'learning_rate' or use the legacy optimizer, e.g., tf.keras.optimizers.legacy.SGD.
Model: "sequential"

Layer (type) Output Shape Param #
=====
conv2d (Conv2D) (None, 32, 32, 32) 896
dropout (Dropout) (None, 32, 32, 32) 0
conv2d_1 (Conv2D) (None, 32, 32, 32) 9248
max_pooling2d (MaxPooling2D) (None, 16, 16, 32) 0
flatten (Flatten) (None, 8192) 0
dense (Dense) (None, 512) 4194816
dropout_1 (Dropout) (None, 512) 0
dense_1 (Dense) (None, 10) 5130
=====
Total params: 4210090 (16.06 MB)
Trainable params: 4210090 (16.06 MB)
Non-trainable params: 0 (0.00 Byte)

None
Epoch 1/5
1563/1563 [=====] - 230s 147ms/step - loss: 1.9127 - accuracy: 0.3121 - val_loss: 1.6984 - val_accuracy: 0.4071
Epoch 2/5
1563/1563 [=====] - 232s 149ms/step - loss: 1.6309 - accuracy: 0.4190 - val_loss: 1.4807 - val_accuracy: 0.4836
Epoch 3/5
1563/1563 [=====] - 243s 155ms/step - loss: 1.4838 - accuracy: 0.4701 - val_loss: 1.4171 - val_accuracy: 0.5074
Epoch 4/5
1563/1563 [=====] - 230s 147ms/step - loss: 1.3967 - accuracy: 0.5024 - val_loss: 1.3227 - val_accuracy: 0.5316
Epoch 5/5
1563/1563 [=====] - 230s 147ms/step - loss: 1.3265 - accuracy: 0.5257 - val_loss: 1.3305 - val_accuracy: 0.5276
Accuracy: 52.76%
```